




(a) Train loss & combined val AUC vs epoch (raw, no smoothing)



(b) xG regression val MAE (best=0.0341 @ ep 24, final=0.0355)



(c) Per-task val-AUC noise (epoch-to-epoch σ , last 10 epochs)



(d) Per-task val AUC vs epoch (8 binary tasks, easy -> hard by max AUC; tight y-range exposes per-task noise; star = best combined epoch; box reports max-skill gap and σ)

